

# A Survey of R-Trees

Aaron Gorenstein  
University of Wisconsin - Madison  
agorenst@cs.wisc.edu

Rebecca Lam  
University of Wisconsin - Madison  
rjlam@cs.wisc.edu

March 9, 2013

## 1 Introduction

What was the motivation for this survey? What was the point of the original paper?

An overview of the rest of the paper here.

## 2 Overview of R-Trees

To solve the problem of performing efficient searches on spatial data, Guttman proposed the R-tree, which inspired a variety of different variations analogous to the family of B-trees. In Section 2.1 we outline the original R-tree paper, and in Section 2.2 we examine the variants and draw appropriate comparisons.

### 2.1 R-Trees

In 1984, Guttman first proposed the idea of modifying the B-tree structure to use minimum bounding rectangles (MBR) as a way to restrict the search space during a lookup for spatial data. This data structure is called the R-tree. R-trees are structured similarly to B-trees except, instead of having separation values in each internal node that divide its subtrees, R-tree internal node entries correspond to MBRs that bound its descendents. For instance, the MBR of a particular node completely overlaps the MBRs of the nodes of its child and its child's children. Like in the B-tree case, nodes correspond to disk pages and leaves point to database objects.

R-trees are bound by two parameters  $m$  and  $M$ , the minimum and maximum number of entries for each node except the root, respectively. An internal node entry is

of the form  $(mbr, p)$ , where  $mbr$  is the MBR containing the MBRs of its descendents and  $p$  is the pointer to its child subtree. The  $mbr$  entry is of the form  $(I_0, I_1, \dots, I_{n-1})$ , where  $n$  is the number of dimensions and  $I_i$  is of form  $[a, b]$ , a closed bounded interval along the  $i$ -th dimension. Similarly, a leaf node entry is of the form  $(mbr, oid)$ , where  $mbr$  is the MBR containing the object, and  $oid$  is the identifier for the object in the database. Finally, the root node must have at least three entries except if it is a leaf.

#### 2.1.1 Search

In order to find all entries contained by a bounding rectangle in the R-tree, the pseudocode of Figure 1 is used.

#### 2.1.2 Insert

Insert.

#### 2.1.3 Delete

Delete

## 2.2 R-Tree Variants

Much like its cousin, the B-tree, the R-tree has a few main variants such as the  $R^+$ -tree and the  $R^*$ -tree, which we discuss in the following sections.

```

function SEARCH( $T, S$ )  $\triangleright$  Return all entries contained
by  $S$  given an R-tree rooted at  $T$ 
    if  $T$  is not a leaf then
        for all  $E$  in  $T$  do
            if  $E.mbr$  overlaps  $S$  then
                SEARCH( $E.p, S$ )
            end if
        end for
    else
        for all  $E$  in  $T$  do
            if  $E.mbr$  overlaps  $S$  then return  $E.oid$ 
            end if
        end for
    end if
end function

```

Figure 1: Pseudocode for searching a R-tree given a search rectangle

#### 2.2.1 R+-Trees

#### 2.2.2 R\*-Trees

### 3 Implementation Challenges

### 4 Database Related Challenges

### 5 Modern Applications

cool topics like parallelism, concurrency, emerging applications

### 6 Conclusion